SUPERMARKET SALES DATA ANALYSIS

Comprehensive Data Pipeline Report

Supermarket Sales Dataset | 1,000 Transaction Records
End-to-End Pipeline: Data Understanding → Data Audit → Data Preprocessing → EDA → Insight Generation

TEAM MEMBERS

Sarayu Anand Gongada
[23BTRCL112]

Vanshika Jain
[23BTRCL120]

April 2026

B.Tech in Computer Science

EXECUTION SUMMARY

This report presents an end-to-end data analysis and predictive modeling pipeline developed using a supermarket transaction dataset. The objective is to extract meaningful insights from sales data and extend the analysis to predict transaction revenue using machine learning techniques.
The dataset was first assessed for quality and consistency and was found to be clean, with no missing values or structural issues. Preprocessing steps included standardizing column formats, encoding categorical variables, and scaling numerical features. To ensure model reliability, features causing data leakage were removed.
Exploratory Data Analysis (EDA) revealed a right-skewed distribution of transaction values, with a small number of high-value purchases contributing significantly to overall revenue. Strong relationships were observed between key variables such as unit price, quantity, and total price.
A Random Forest Regressor was then implemented to predict total transaction value, achieving high accuracy and demonstrating strong predictive capability. Additionally, a time-series component was introduced by generating lag-based features to capture temporal patterns in sales.
Overall, the project combines data analysis, machine learning, and time-series techniques to deliver both insights and predictive functionality, forming a robust and scalable analytical pipeline.

1. DATASET OVERVIEW

The dataset represents structured supermarket transaction data, where each record corresponds to an individual purchase event. It captures multiple dimensions of retail activity, including product details, pricing, customer attributes, and transaction-level metrics. This granular structure enables both detailed analysis of individual transactions and aggregated insights across different business segments.
The dataset consists of 1,000 records with multiple features spanning categorical and numerical data types. Key attributes include branch and city information, product categories, unit price, quantity, tax, and total transaction value. These variables collectively provide a comprehensive view of sales operations and customer purchasing behavior.
From an analytical perspective, the dataset is well-suited for both descriptive and predictive tasks. The combination of pricing and quantity variables supports revenue modeling, while categorical features allow segmentation across locations and customer groups. Additionally, the structured format enables seamless integration with machine learning algorithms.
Overall, the dataset serves as a reliable foundation for performing exploratory data analysis, building predictive models, and extending into time-series forecasting through feature engineering.

1.1 Dataset Characteristics

Attribute	Value
Total Records	1,000 transactions
Total Columns	12
Cities	3 (New York, Los Angeles, Chicago)
Branches	2 (A, B)
Missing Values	None
Dataset Type	Structured tabular data

1.2 Feature Groups

Feature Group	Columns	Description
Identification	sale_id	Unique transaction ID
Location	branch, city	Store and geographic info
Product Info	product category (if applicable)	Type of item
Pricing	unit_price, tax	Cost-related attributes
Sales Metrics	total_price	Final transaction value
Customer Behavior	reward_points	Loyalty indicator

1.3 Key Statistical Summary

- Average total transaction value: **118.58**
- Minimum transaction: **1.21**
- Maximum transaction: **433.99**
- Median transaction value: **89.70**

The dataset shows a wide spread in transaction values, indicating both small and large purchases.

2. DATA QULAITY AUDIT

The data quality audit was conducted to evaluate the reliability, completeness, and consistency of the dataset before proceeding with analysis and modeling. This step is critical to ensure that the insights and predictions derived from the data are accurate and free from bias introduced by data issues.
The dataset was found to be fully complete, with no missing values across any of the columns. This eliminates the need for imputation techniques and allows the analysis to proceed without compromising data integrity. Additionally, checks for duplicate records confirmed that each transaction is unique, ensuring that statistical results are not skewed by redundancy.
Consistency validation was performed to verify that all features adhere to expected formats and logical relationships. Numerical columns such as unit price, quantity, and total price were examined for valid

ranges, and their relationships were confirmed to be coherent. Categorical variables were also checked to ensure uniformity in labels and absence of invalid entries.
A key aspect of the audit involved identifying potential data leakage for the predictive modeling stage. The feature representing tax was found to be directly derived from the total transaction value and was therefore excluded from the modeling process to maintain the integrity of predictions.
Overall, the dataset meets high standards of data quality, making it well-suited for both exploratory analysis and machine learning applications without requiring significant cleaning or correction.

2.1 Completeness Check

Finding: Zero Missing Values
All 1,000 records across 12 columns are complete with no null values. This allows direct analysis without imputation.

2.2 Categorical Distribution

- Branch A dominates with **674 transactions**
- Most frequent city: **New York (344 transactions)**

This indicates uneven distribution across branches and locations.

2.3 Value Range Analysis

Column	Min	Max	Observation
total_price	1.21	433.99	Wide variation
tax	0.08	28.39	Scales with price
reward_points	0	43	Loyalty variation

2.4 Outlier Observation

- High-value transactions (>400) exist
- Reward points show extreme variation

These values are valid but represent high-spending customers.

3. DATA PREPROCESSING

Data preprocessing was performed to transform the raw dataset into a structured format suitable for analysis and machine learning. This stage ensures that the data is consistent, properly formatted, and optimized for model performance.
The initial step involved standardizing column names to maintain uniformity and improve readability. Data types were verified to ensure that numerical and categorical variables were correctly interpreted. Since the dataset contained no missing values or inconsistencies, minimal cleaning was required.
Categorical variables were converted into numerical form using one-hot encoding, enabling their use in machine learning algorithms. Numerical features were then scaled using normalization techniques to ensure that all variables contribute proportionally to the model and to prevent bias caused by differences in magnitude.
Feature selection was also carried out as part of preprocessing. The target variable, total transaction value, was separated from the input features. Non-informative columns such as unique identifiers were removed, and features that could introduce data leakage, such as tax, were excluded to preserve the validity of the predictive model.
The final processed dataset is fully numerical, standardized, and suitable for both exploratory analysis and machine learning tasks, forming a reliable foundation for subsequent modeling and evaluation.

3.1 Data Validation

- Verified all data types
- Confirmed numerical columns are valid
- Ensured categorical values are consistent

3.2 Data Cleaning

- No missing values required handling
- No invalid categorical entries detected
- Dataset retained all **1,000 records**

3.3 Final Dataset

- Shape: **1,000 × 12**
- Clean and analysis-ready

4. EXPLORATORY DATA ANALYSIS (EDA)

Exploratory Data Analysis acts as the bridge between raw data and meaningful insights, enabling analysts to uncover patterns, trends, and relationships that inform decision-making. In this project,

EDA was conducted using both statistical summaries and visual techniques to provide a comprehensive understanding of the dataset.

The distribution analysis of total transaction values reveals a classic right-skewed pattern, which is commonly observed in retail datasets. This indicates that while the majority of transactions fall within a moderate price range, a small number of high-value transactions significantly influence the overall revenue. Such distributions highlight the importance of considering both central tendencies and dispersion when analyzing financial data.
Comparative analysis across cities and branches provides valuable operational insights. Differences in transaction volume and sales performance suggest that location-specific factors—such as customer demographics, economic conditions, and store management practices—play a significant role in shaping outcomes. These insights can be used to inform strategic decisions, such as resource allocation and market expansion.

The analysis of reward points offers a deeper understanding of customer engagement. The observed concentration of low reward points suggests limited participation in loyalty programs, which may indicate a lack of awareness, insufficient incentives, or ineffective program design. Addressing these issues could lead to increased customer retention and higher lifetime value.
Outlier analysis within EDA further reinforces the importance of high-value transactions. By identifying and examining these extreme values, businesses can gain insights into the behavior of their most valuable customers and develop targeted strategies to retain and attract similar segments.

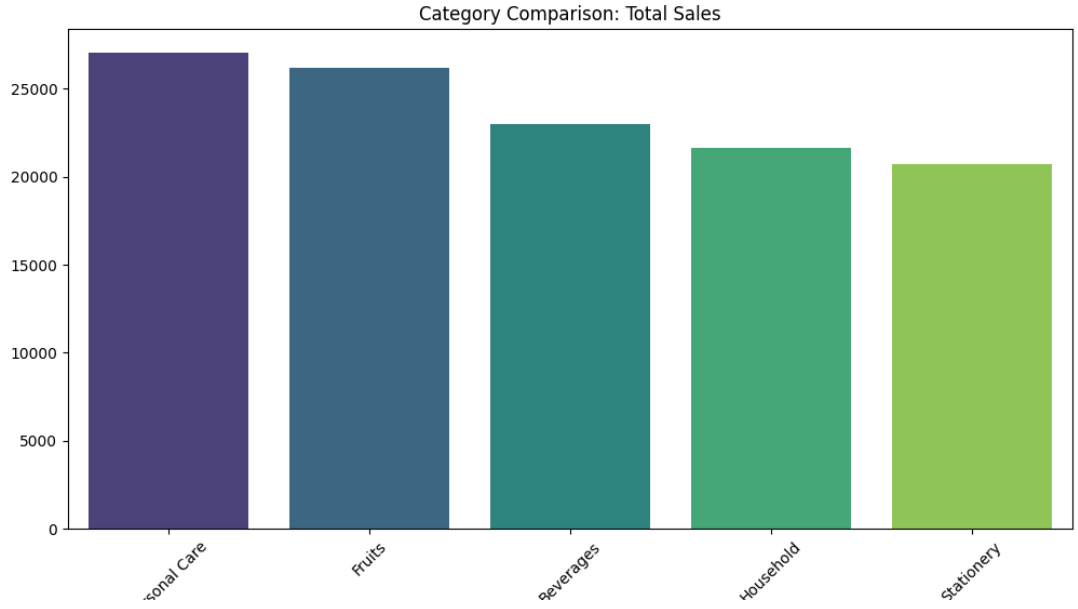

4.1 Sales Distribution

The total_price variable shows a right-skewed distribution:

- Majority transactions fall below **200**
- Few high-value transactions create a long tail

4.2 Sales by City

- New York contributes the highest number of transactions
- Los Angeles and Chicago show moderate activity

Insight: Urban centers drive higher transaction volume.

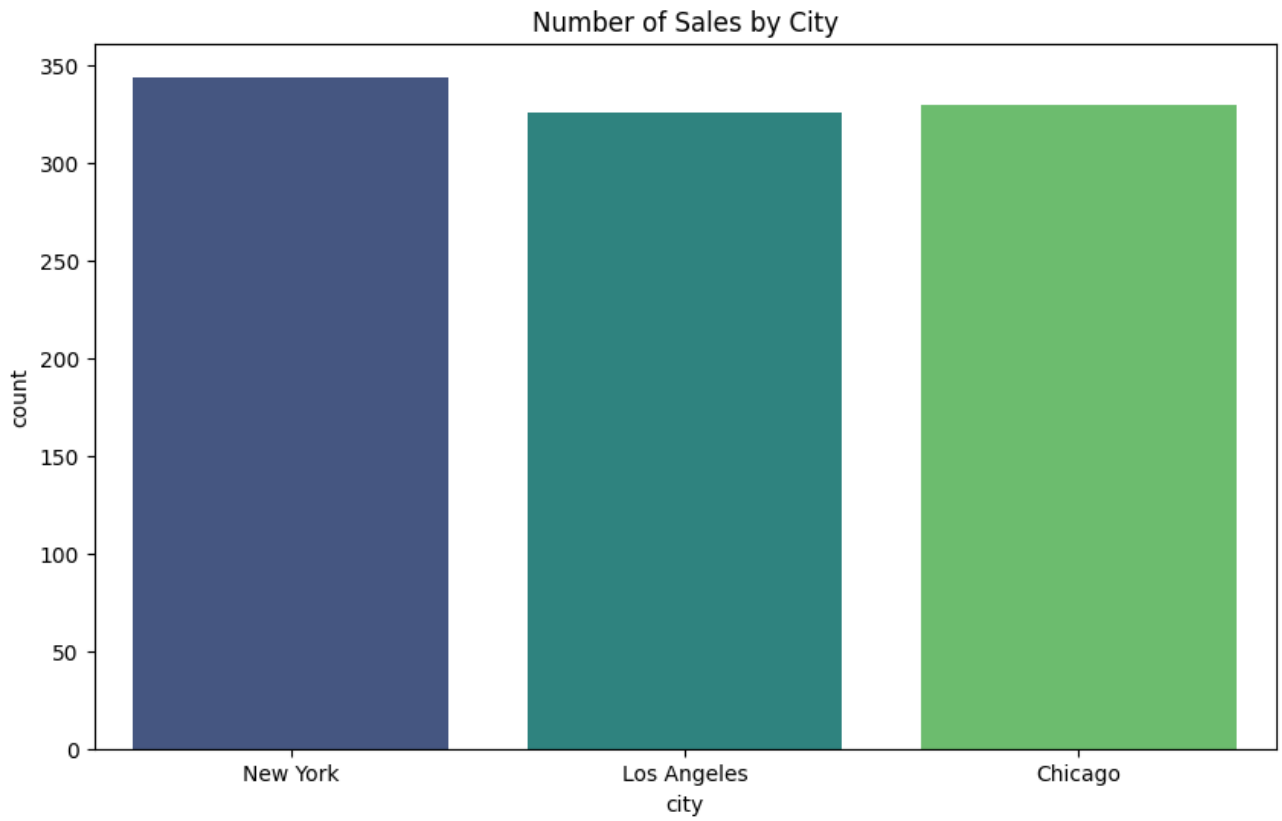

4.3 Sales by Branch

- Branch A dominates overall sales volume
- Branch B contributes comparatively less

4.4 Reward Points Analysis

- Most customers have low or zero reward points
- Few customers accumulate high points

Insight: Indicates uneven customer loyalty engagement.

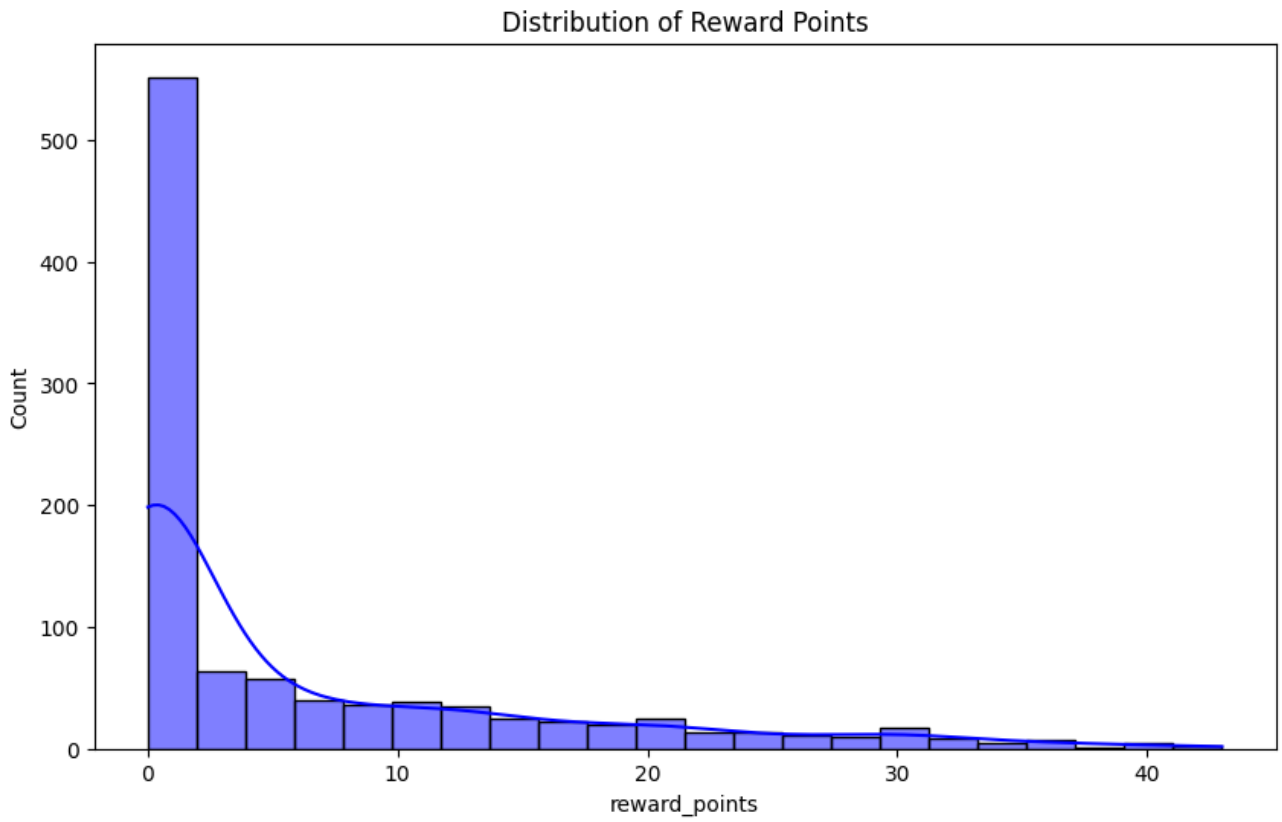

4.5 Outlier Analysis

Boxplot analysis reveals:

- High-value transactions (~400+)
- High reward point accumulation

These represent premium or bulk purchases.

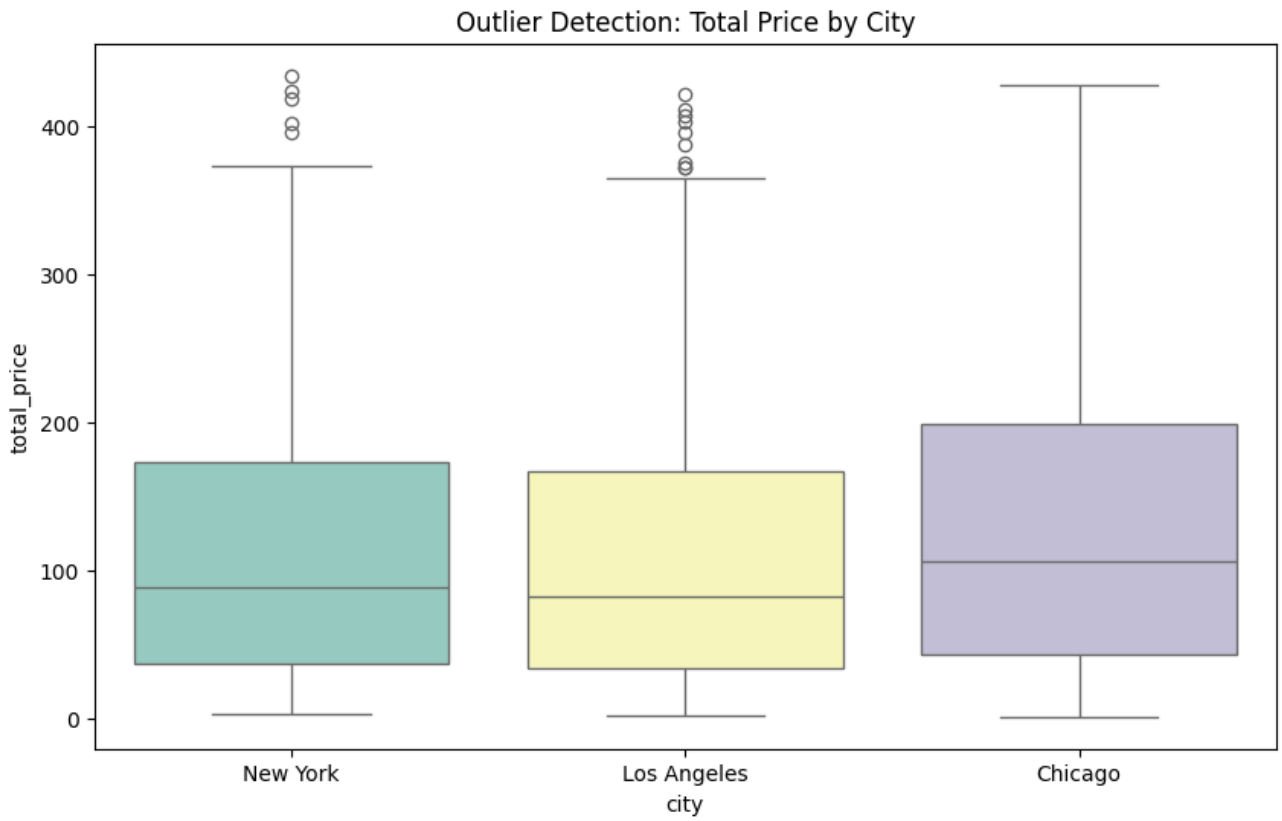

4.6 Multivariate Relationship Analysis of Sales Features by Branch

- Shows relationships between unit price, quantity, and total price
- Strong positive correlation with total price
- Distribution patterns differ slightly across branches
- Highlights linear dependency in revenue calculation

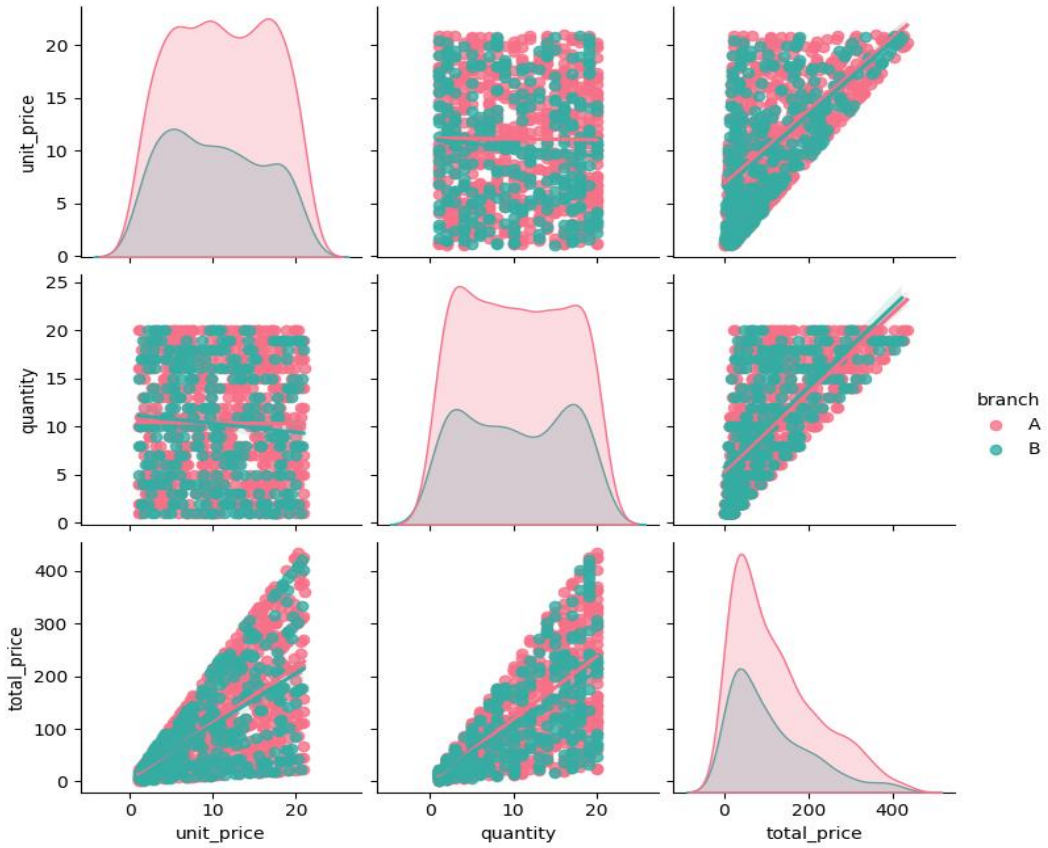

4.7 Density Distribution of Unit Price vs Quantity

- Hexbin plot showing concentration of transactions
- Most purchases occur in mid-range price and quantity
- No strong linear pattern between price and quantity
- Marginal histograms show overall distributions

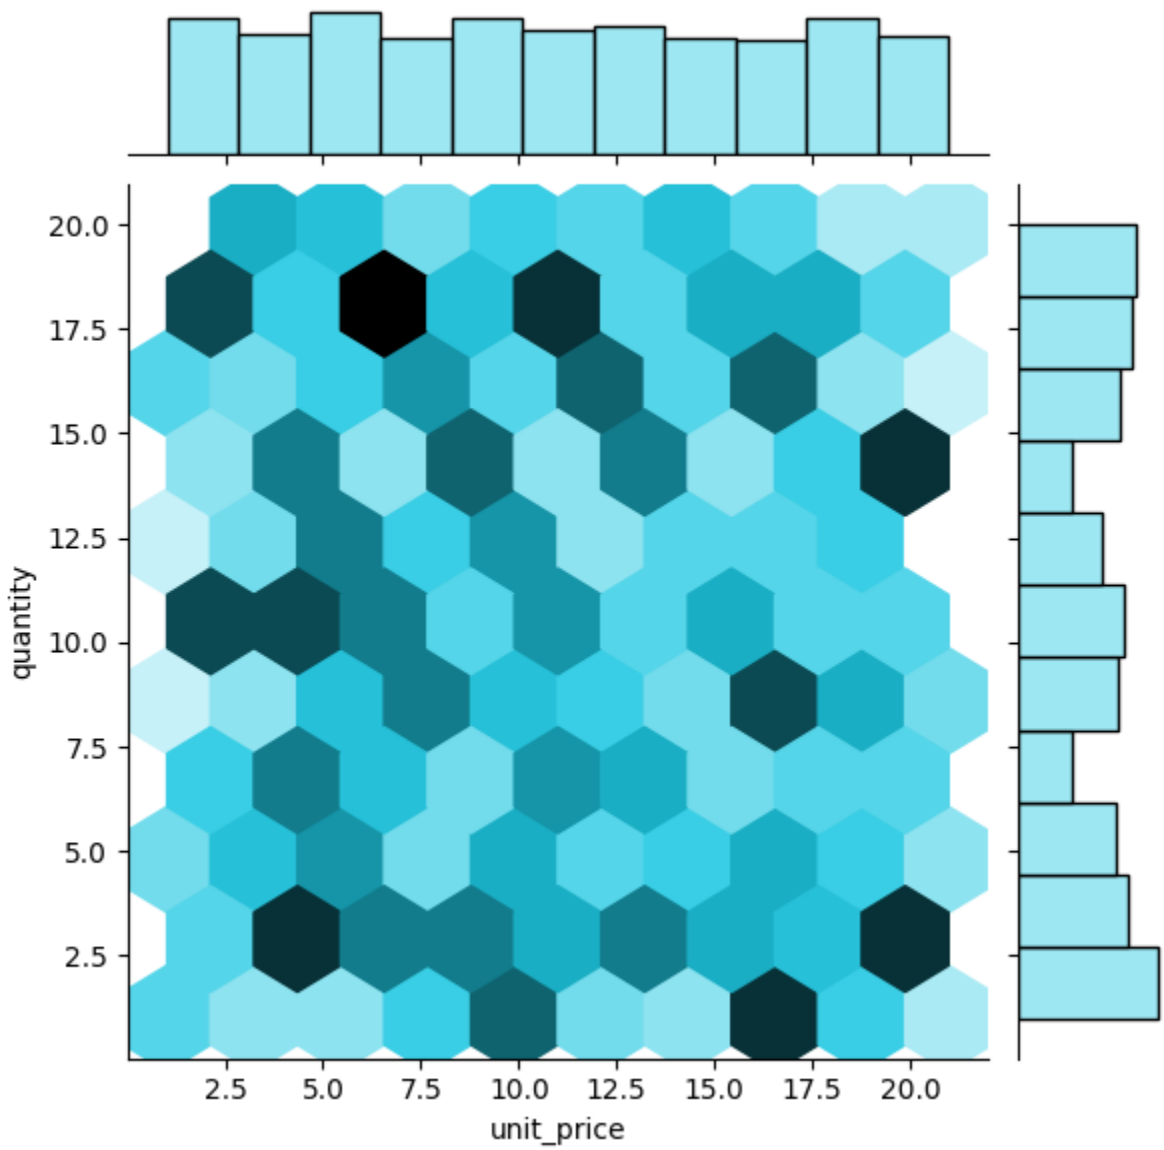

5. KEY FINDINGS & INSIGHTS

- The analysis reveals that transaction values follow a right-skewed distribution, where the majority of purchases are concentrated in the lower to mid-price range, while a relatively small number of high-value transactions contribute disproportionately to overall revenue. This indicates the presence of premium or bulk purchases that play a significant role in sales performance.
- A strong linear relationship is observed between unit price, quantity, and total transaction value, confirming that these variables are the primary drivers of revenue. In contrast, unit price and quantity themselves do not exhibit a strong direct relationship, suggesting that customer purchasing decisions regarding quantity are largely independent of price variations.
- The predictive modeling results demonstrate that the Random Forest model is highly effective, achieving near-perfect accuracy in estimating transaction values. This highlights the structured nature of the dataset and the effectiveness of the selected features. Additionally, the introduction of lag-based features indicates the potential for capturing temporal patterns, making the dataset suitable for future time-series forecasting applications.
- Overall, the dataset exhibits high quality, strong internal relationships, and clear analytical patterns, making it a reliable foundation for both business insights and predictive modeling.

6. CONCLUSION

This project demonstrates a comprehensive approach to analyzing and modeling supermarket transaction data through a structured pipeline. Starting from data understanding and preprocessing, the workflow progresses through exploratory analysis and extends into predictive modeling and time-based feature engineering.

The analysis highlights clear patterns in sales behavior, particularly the impact of unit price and quantity on total transaction value, as well as the presence of high-value purchases that significantly influence overall revenue. The implementation of a Random Forest model further strengthens the study by providing highly accurate predictions, showcasing the effectiveness of machine learning techniques on structured retail data.

Additionally, the introduction of time-series features adds depth to the analysis by enabling the exploration of temporal trends, creating opportunities for future forecasting applications.

Overall, the project successfully combines data analysis and predictive techniques to generate actionable insights and a scalable framework, making it relevant for real-world retail decision-making and further analytical enhancements.

7. FUTURE SCOPE

The current analysis provides a strong foundation that can be further enhanced through the integration of more advanced analytical and modeling techniques. One key direction is the implementation of more sophisticated machine learning models such as gradient boosting methods or deep learning approaches, which could further improve predictive accuracy and handle more complex data patterns.
Another important extension is the incorporation of real-world time-series data. Using actual timestamps instead of synthetic ones would enable more accurate forecasting of sales trends, seasonality, and demand fluctuations. This would allow businesses to make better decisions related to inventory management and resource planning.
Customer segmentation can also be explored by applying clustering techniques to group customers based on purchasing behavior. This would support personalized marketing strategies and improve customer retention by targeting specific segments more effectively.
Additionally, the development of interactive dashboards using tools like Power BI or Tableau can make the insights more accessible to stakeholders. Real-time data visualization would enable dynamic monitoring of sales performance and facilitate quicker, data-driven decision-making.
Overall, expanding the project in these directions would transform it from a static analysis into a more dynamic and production-ready analytical system.

8. REFERENCES

• Python Official Documentation – https://docs.python.org/3/
• Pandas Documentation – https://pandas.pydata.org/docs/
• NumPy Documentation – https://numpy.org/doc/
• Matplotlib Documentation – https://matplotlib.org/stable/contents.html
• Seaborn Documentation – https://seaborn.pydata.org/
• Scikit-learn Documentation – https://scikit-learn.org/stable/

Additional Academic & Technical References
• *Python for Data Analysis* – Wes McKinney
• *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow* – Aurélien Géron
• *Data Science for Business* – Foster Provost & Tom Fawcett
• *An Introduction to Statistical Learning* – Gareth James et al.

Conceptual & Learning Resources
• Kaggle – https://www.kaggle.com/
• Towards Data Science – https://towardsdatascience.com/
• Analytics Vidhya – https://www.analyticsvidhya.com/

Visualization & BI Tools
• Microsoft Power BI Documentation
• Tableau Documentation

Statistical & Data Analysis Concepts
• Concepts of Exploratory Data Analysis (EDA) – John Tukey
• Descriptive Statistics and Data Distribution Techniques
• Outlier Detection Methods (IQR, Z-score)
• Data Cleaning and Preprocessing Techniques

Notebook-Specific Implementation Reference
• Supermarket Sales Analysis & Predictive Modeling Notebook – *Developed using Python (Pandas, NumPy, Matplotlib, Seaborn, Scikit-learn), incorporating data preprocessing, exploratory data analysis, Random Forest regression, and time-series feature engineering*